



Approximate Pattern Matching

Riade Benbaki, Ali Mammadov

INF560 Project defense

Plan

1. Approximate pattern matching problem
2. Parallelism approaches
 - MPI
 - OpenMP
 - CUDA
 - Mixing them all together
3. Results and experiments



1 - Approximate pattern matching problem

- We consider each pattern independently
- For each pattern, we check if there is a match between the pattern and the string and the current starting position
- We count the total number of matches obtained

What can we parallelize ?

- Over the different patterns
- Over each starting position in the file
- Not over the distance computation
 - Based on Levenshtein distance -> Hard to parallelize



MPI

- Split the file evenly over all MPI ranks (data parallelism)
- Every rank received it's own part + some shadow cells necessary for the last computations
- An MPI reduction at the end to sum up the results of all ranks



OpenMP

- For every MPI rank, we use OpenMP threads to parallelize the loop over all starting positions
- Every threads needs to read the string, but not modify it
- Every thread needs a private array for computation



CUDA

- The deepest level that can be parallelized is the previous
- We therefore use CUDA with OpenMP, each of them processing a part of the chunk of the current MPI rank
- One cuda thread needs access to
 - String buffer -> size of `buf_size = n_bytes/size`
 - Pattern -> size of at most `max_pat`
 - An array column (private to each thread) -> `n_threads` array of ints
 - A results array -> `n_threads` array of ints
- All this for each MPI process -> multiple ranks may share same GPU
- Worst case scenario : all ranks share the same GPU

$$[nth_{max}(max_pat + 2) \times sizeof(int) + buf_size \times sizeof(char)] size \leq freeGpUMemory()$$

- Use a ratio value to divide the work between CUDA and OpenMP

Parallel Algorithm

- kernelCall() - calls kernel function
 - Asynchronous
- finalcudaCall() - Synchronizes kernel
 - Returns results
- Mpi_Reduce() - For getting the
 - results from each process

Result: Number of the approximate patterns

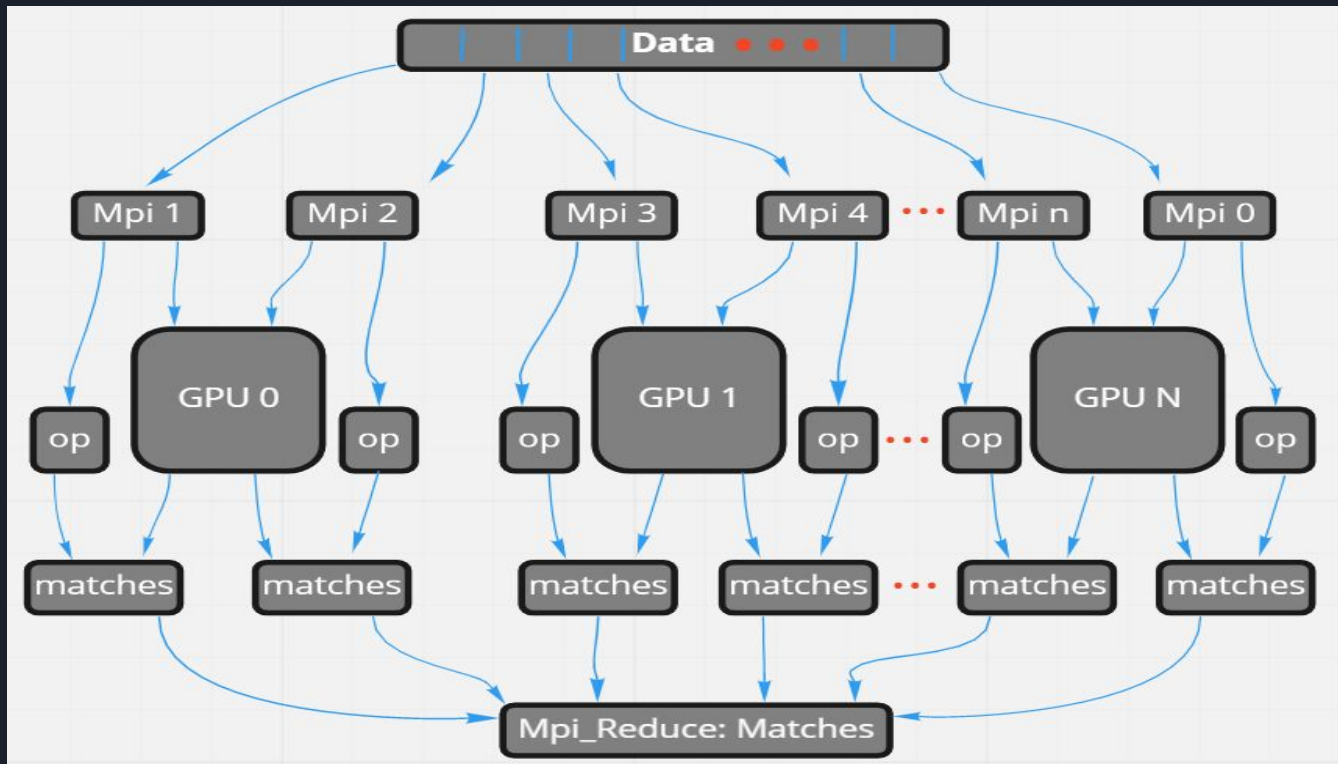
Init:

```
Allocate tables;
Send tables to devices;
for each p in length patterns do
    send pattern to devices
    kernelCall();
    #pragma omp parrallel
    OpenMp initialization;
    #pragma omp for reduction ompMatches
    for each j in OpenMp part do
        | ompMatches += (levenshtein() < aproxFactor) ;
    end
    #pragma omp single
    Synchronize kernel -> finalcudaCall( results );

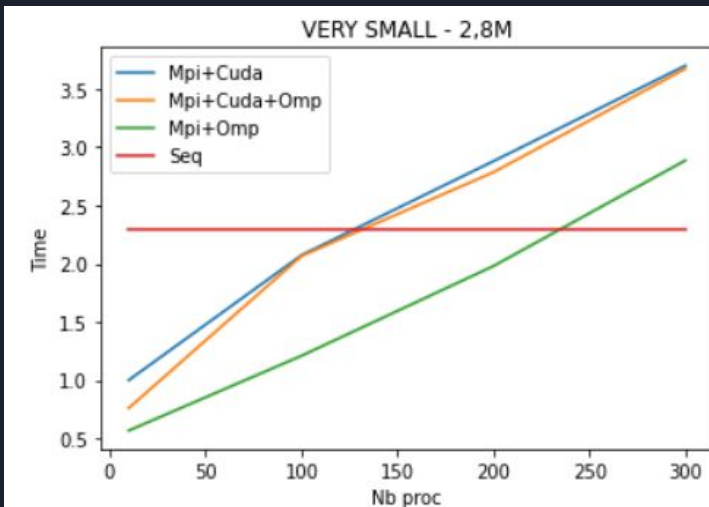
    #pragma omp for reduction cudaMatches
    cudaMatches = sum(results);

    matches[p] = cudaMatches + ompMatches;
    free( some memory );
end
MpiReduce(matches[p]);
free( some memory );
```

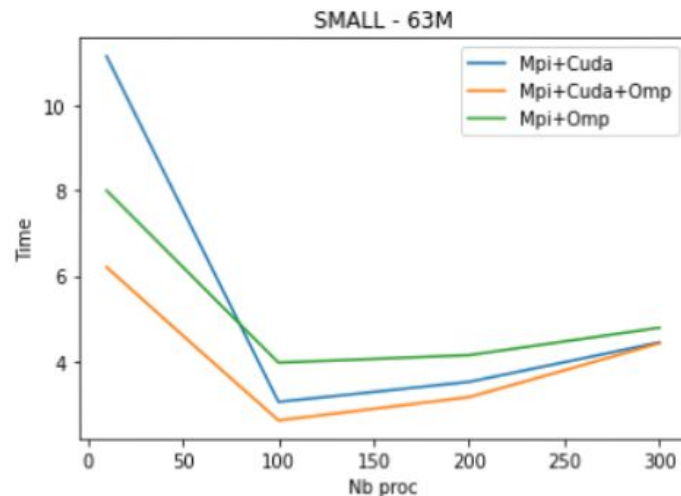
Parallel Algorithm



Experimental evaluations - small files

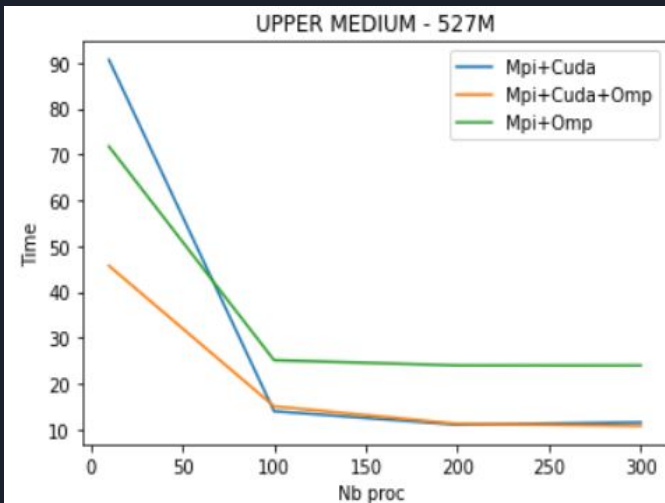


Minimum speedup = 0.6186212674082705
Maximum speedup = 4.046322719955613




Minimum speedup = 3.570232421317228
Maximum speedup = 15.331004418483946

Experimental evaluations - large files



Minimum speedup = 4.571884137171331
Maximum speedup = 38.61584246698981



Minimum speedup = 3.732405832628397
Maximum speedup = 34.67503064440221

Thank you for your attention